

Exkurs: linear Algebra

Matrix: $\begin{bmatrix} 5 & 12 & 6 \\ 3 & 0 & 14 \end{bmatrix} = A_{2x3}$ $a_{12} = 12$ 2 Dimensions

Scalars: $[15] = A_{1x1}$ 0 Dimensions

Vector: $\begin{bmatrix} 5 \\ -2 \\ 4 \end{bmatrix} = V_{3x1}$ 1 Dimensions

Transpose:

$$X = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \quad X^T = [1 \ 2 \ 3] \quad (X^T)^T = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

Dot Product:

$$m \times n \cdot n \times k = m \times k$$

$$3 \times 1 \cdot 1 \times 4 = 3 \times 4$$

NEW AIRING: Udemy Deep Learning

Training an Algorithm:

- Data
- Model
- Objective Function \rightarrow minimize error
- Optimization Algorithm

Linear Model: $f(x) = w \cdot x + b$

↑ ↑
 Input coefficient Intercept
 (weight) (Bias)

Example: Houseprice = sqft(x) * weight - bias *

Input Model is determined by W & B

$$Y = XW + B$$

= $n \times m$ * $n \times m$ = $1 \times m$
 Obs Input In Out

n \times m 1 \times m
 Observation Out
 Output var

Objective Functions

def: Measure used to evaluate how well the Output matches the desired correct Values

Regression: L2-norm = $\sum_i (y_i - t_i)^2$

Lower Loss
 =
 Lower Error

output target

Classification: Cross-Entropy = $L(y, t) = -\sum_i t_i \ln y_i$

Inputs (Cats, dogs)

Optimization Algorithms

Gradient descent:

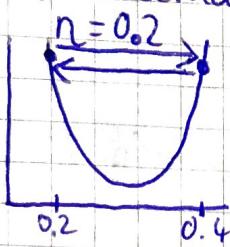
RL-Example 8 $f'(x) = 10x + 3$

$$\begin{aligned} x_0 &= 4 \\ x_1 &= 4 - n \cdot 43 \end{aligned}$$

x_0
 x_{not}

$n(\eta)$ → Learning rate

if learningrate
too big we will
start oscillate



the minimum is reached

when η

$$x_{i+1} = x_i - n f'(x_i)$$

$$x_{i+1} = x_i \rightarrow \text{no more updates}$$

TensorFlow 2.0



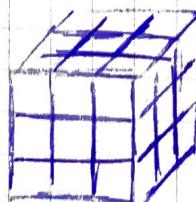
Skalar Tensor
Tensor Rank1



Vector Tensor
Tensor Rank2



Matrix Tensor
Tensor Rank3



Tensor

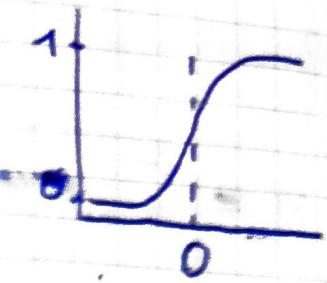
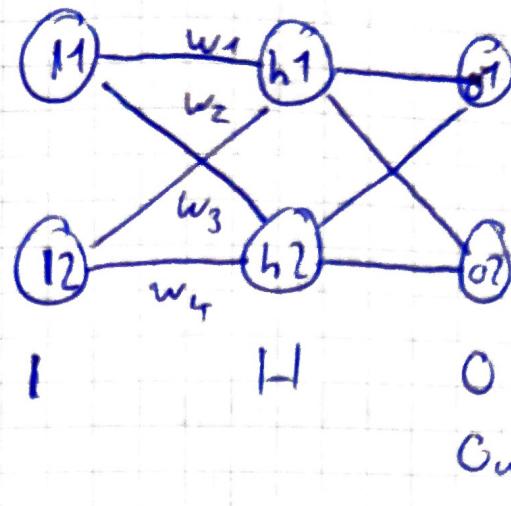
data → preprocess → save in .npz

Multidimensional
Files

easier and faster
to work with

Epoch = iteration over the full dataset

Forward Pass



$$net_{h_1} = w_1 * i_1 + w_2 * i_2 + b_1 * 1$$

Implementation:

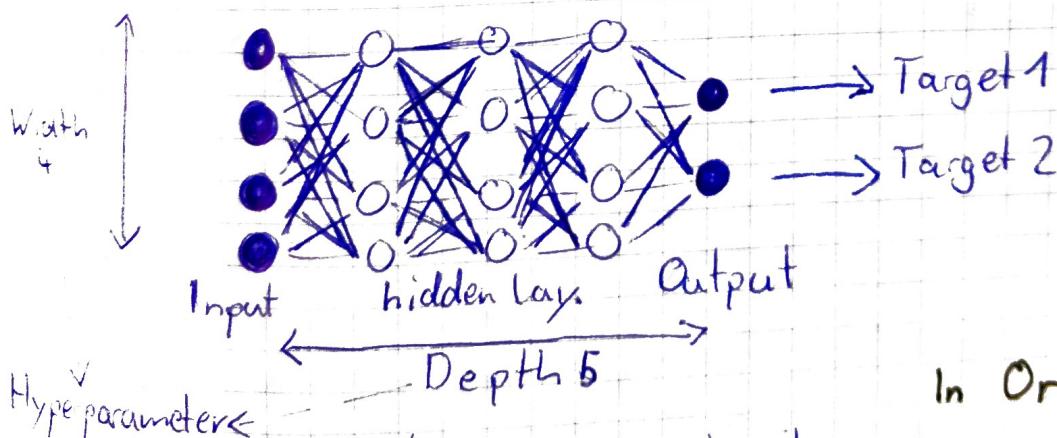
Scaling is very important!!!

Front and Elbow needs to be scaled clearly

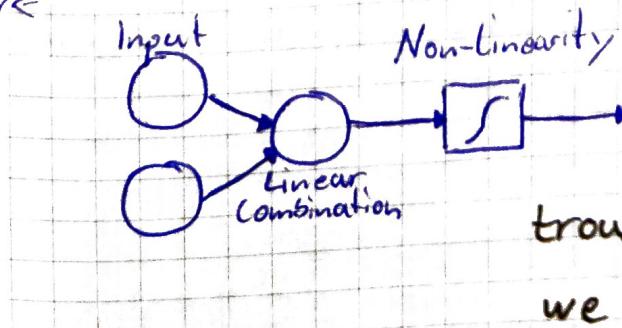
Maybe write own Trainer, Keras might be enough

Lecture by Jeroen

What is a deep Net?



Hyperparameters



In Order to have deep nets and find complex relationships through arbitrary functions, we NEED non-linearities

non-linearities are called ACTIVATION FUNCTION

↳ also: Transfer functions

The last activation function of a dNN is often

$$\text{a Softmax} \quad \frac{e^{a_i}}{\sum e^{a_j}} = \frac{E}{n}$$

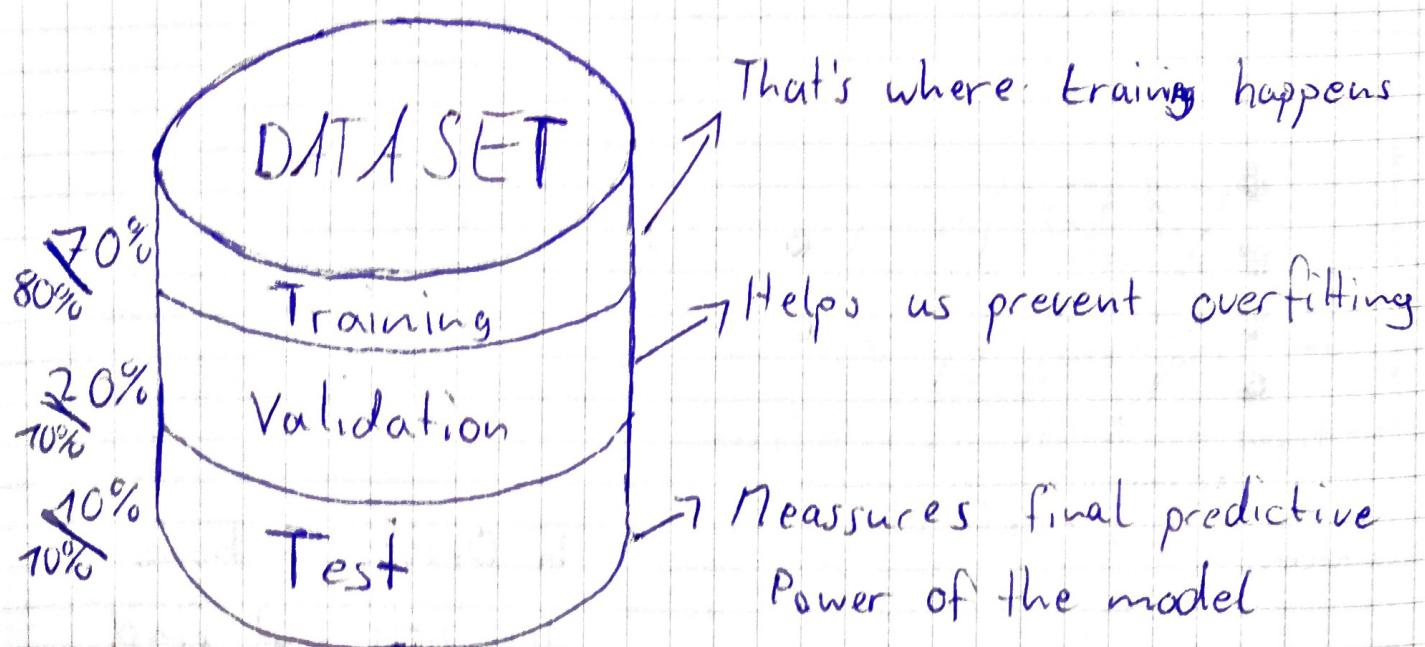
↳ This gives back PROBABILITIES that add up to 1. Which makes the softmax very convenient for the last hidden Layer

What's Overfitting:

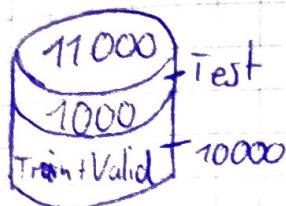
Our training has focused on the particular training set so much, it has "missed the point"

What's Underfitting:

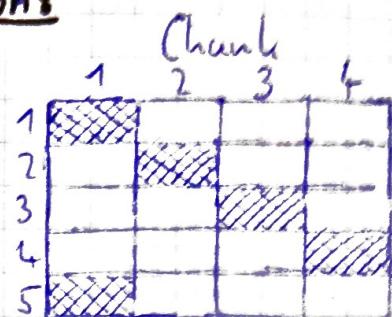
The model has not captured the underlying logic of the data.



N-fold cross-validation:



1
Epoch



Train
Valid

Pros:

utilized more data



We have a Model

tradeoff

Red flag

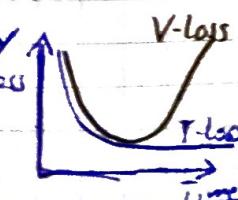
Cons:

Possibly, overfitted a bit

Only used if not enough data!!!

Early Stopping -

	<u>Pros</u>	<u>Cons</u>
1 Preset number of Epochs	- Eventually, solves the Problem	- No guarantee to reach min - Might not minimize - Naive
2 Stop when updates become too small	- We're sure the loss is minimized - Saves computing power	- Might still overfit
3 Validation set Strategy	- We're sure the loss is minimized - Saves computing Power - Prevents overfitting	- Might iterate through validation set widely
4 best parts of 1, 2, and 3	all of them	NONE



PREPROCESSING

Why? :

- Compatibility
- Orders of magnitude
- Generalization

	bread	yogurt	muffin
1	1	0	0
2	0	1	0
3	0	0	1

1. Standardization:

$$\text{standardized variable} = \frac{x - \mu}{\sigma}$$

original variable
 mean of original Variable
 standard deviation of original variable

2. Normalization:

5	100
120	-10
-1	0

normalize
using L2-norm

0.042	0.995
0.999	-0.1
0.003	0

std =
standard deviation

3. PCA:

principal components analysis

x_r y_r z_r rotation where $\text{mean} = 0$
 $\text{std} = 1$

4. Whitening:

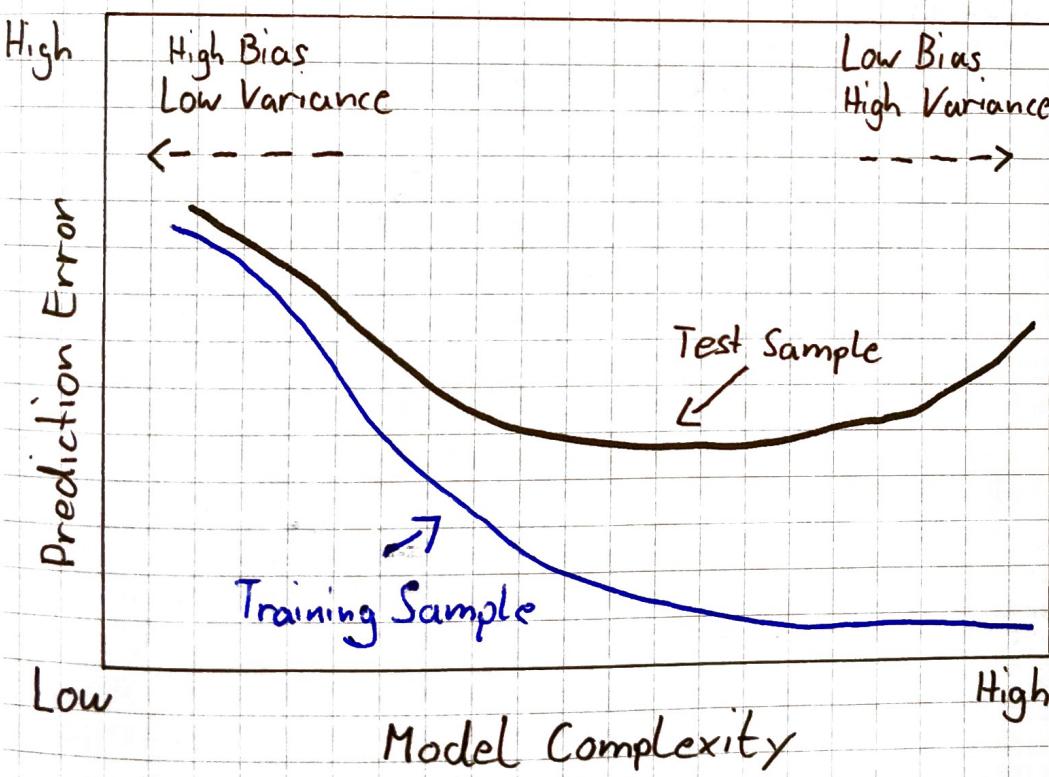
⋮

Week 4: Evaluation Methodology

Regression: Estimate the relationship between a target variable and the Features (House vs. Room)

Classification: predict the most likely class given the features of a datapoint

Validation



Evaluation

- Machine Learning is a Black-box
 - We make many assumptions
 - About the Model and its type
 - About interpretation of Data
 - Validation: did we learn the model correct?
 - Evaluation: did we learn the correct model?
- Rules of thumb for sound Machine Learning
 - Data: Check the quality of your data
 - Model: Law of Parsimony
 - Validation
 - Evaluation
 - measure effectiveness
 - Compare against baseline (sanity check)

LAW OF
PARSIMONY

The principle that
the simplest
Explanation of an
event or observation
is the preferred
explanation

- Effectiveness
 - the degree to which a learned model is successful in producing a desired result

- Efficiency
 - Computation time
 - Memory use

Evaluation Metrics

- Regression/continuous scale
 - Mean Squared Error

$$\frac{1}{m} \sum_{i=1}^m (\hat{y} - y)^2$$

- Reduces variance to data points
- Distant points are more important

- Root Mean Squared Error

$$\sqrt{\frac{1}{m} \sum_{i=1}^m (\hat{y} - y)^2}$$

- More explainable than MSE because it is on the same scale as the target Variable

- Mean Absolute Error

$$\frac{1}{m} \sum_{i=1}^m |\hat{y} - y|$$

- Is more tolerant to outliers

- Regression/continuous scale

$$R^2 = 1 - \frac{\sum_{i=1}^m (\hat{y} - y)^2}{\sum_{i=1}^m (y - \bar{y})^2}$$

Percentage of variance in the target variable that is explained by the features